

yzx0bk4lf

October 17, 2024

```
[1]: #Aim : To perform and find the accuracy of Logistic Regression
```

```
[2]: #Name: Ritika Rajesh Junekar  
#Roll no. :30  
#Sub : ET1  
#section:C
```

```
[1]: import pandas as pd  
import matplotlib.pyplot as plt  
import numpy as np  
import seaborn as sns  
from sklearn.model_selection import train_test_split  
import warnings  
warnings.filterwarnings('ignore')
```

```
[2]: import os
```

```
[3]: os.getcwd()
```

```
[3]: 'C:\\Users\\USER'
```

```
[7]: os.chdir("C:\\Users\\USER\\Desktop")
```

```
[9]: df=pd.read_csv("framingham.csv")
```

```
[11]: #The "Framingham" heart disease dataset includes over 4,240 records, 15  
      ↪ attributes.  
      #The goal of the dataset is to predict whether the patient has 10-year risk of  
      ↪ future (CHD) coronary heart disease
```

```
[13]: df.head()
```

```
[13]:   male  age  education  currentSmoker  cigsPerDay  BPMeds  prevalentStroke  \  
0      1   39         4.0              0          0.0      0.0              0  
1      0   46         2.0              0          0.0      0.0              0  
2      1   48         1.0              1         20.0      0.0              0  
3      0   61         3.0              1         30.0      0.0              0
```

```
4      0      46      3.0      1      23.0      0.0      0
```

```
prevalentHyp  diabetes  totChol  sysBP  diaBP  BMI  heartRate  glucose \
0              0         0    195.0  106.0   70.0   26.97      80.0    77.0
1              0         0    250.0  121.0   81.0   28.73      95.0    76.0
2              0         0    245.0  127.5   80.0   25.34      75.0    70.0
3              1         0    225.0  150.0   95.0   28.58      65.0   103.0
4              0         0    285.0  130.0   84.0   23.10      85.0    85.0
```

```
TenYearCHD
0          0
1          0
2          0
3          1
4          0
```

```
[15]: df.describe()
```

```
[15]:      male      age  education  currentSmoker  cigsPerDay \
count  4238.000000  4238.000000  4133.000000    4238.000000  4209.000000
mean      0.429212    49.584946    1.978950      0.494101    9.003089
std      0.495022    8.572160    1.019791      0.500024   11.920094
min      0.000000   32.000000    1.000000      0.000000    0.000000
25%      0.000000   42.000000    1.000000      0.000000    0.000000
50%      0.000000   49.000000    2.000000      0.000000    0.000000
75%      1.000000   56.000000    3.000000      1.000000   20.000000
max      1.000000   70.000000    4.000000      1.000000   70.000000
```

```
      BPMeds  prevalentStroke  prevalentHyp  diabetes  totChol \
count  4185.000000    4238.000000    4238.000000  4238.000000  4188.000000
mean      0.029630      0.005899      0.310524      0.025720   236.721585
std      0.169584      0.076587      0.462763      0.158316   44.590334
min      0.000000      0.000000      0.000000      0.000000   107.000000
25%      0.000000      0.000000      0.000000      0.000000   206.000000
50%      0.000000      0.000000      0.000000      0.000000   234.000000
75%      0.000000      0.000000      1.000000      0.000000   263.000000
max      1.000000      1.000000      1.000000      1.000000   696.000000
```

```
      sysBP      diaBP      BMI      heartRate      glucose \
count  4238.000000  4238.000000  4219.000000  4237.000000  3850.000000
mean    132.352407    82.893464    25.802008    75.878924    81.966753
std     22.038097    11.910850     4.080111    12.026596    23.959998
min     83.500000    48.000000    15.540000    44.000000    40.000000
25%    117.000000    75.000000    23.070000    68.000000    71.000000
50%    128.000000    82.000000    25.400000    75.000000    78.000000
75%    144.000000    89.875000    28.040000    83.000000    87.000000
max    295.000000   142.500000    56.800000   143.000000   394.000000
```

```

      TenYearCHD
count  4238.000000
mean    0.151958
std     0.359023
min     0.000000
25%     0.000000
50%     0.000000
75%     0.000000
max     1.000000

```

```
[17]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4238 entries, 0 to 4237
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   male                  4238 non-null  int64
1   age                   4238 non-null  int64
2   education             4133 non-null  float64
3   currentSmoker         4238 non-null  int64
4   cigsPerDay            4209 non-null  float64
5   BPMeds               4185 non-null  float64
6   prevalentStroke       4238 non-null  int64
7   prevalentHyp         4238 non-null  int64
8   diabetes              4238 non-null  int64
9   totChol              4188 non-null  float64
10  sysBP                4238 non-null  float64
11  diaBP                4238 non-null  float64
12  BMI                  4219 non-null  float64
13  heartRate            4237 non-null  float64
14  glucose              3850 non-null  float64
15  TenYearCHD           4238 non-null  int64
dtypes: float64(9), int64(7)
memory usage: 529.9 KB

```

```
[19]: df.isna().sum()
```

```

[19]: male                0
      age                0
      education          105
      currentSmoker      0
      cigsPerDay         29
      BPMeds            53
      prevalentStroke    0
      prevalentHyp      0

```

```

diabetes      0
totChol      50
sysBP        0
diaBP        0
BMI          19
heartRate     1
glucose      388
TenYearCHD    0
dtype: int64

```

```

[21]: #Since, only a few rows have null values in them, we are only removing those
      ↪rows from the dataset.
      #df = df.dropna(subset=['heartRate', 'BMI', 'cigsPerDay', 'totChol', 'BPMeds'])

```

```

[23]: df

```

```

[23]:      male  age  education  currentSmoker  cigsPerDay  BPMeds  \
0         1   39         4.0              0         0.0      0.0
1         0   46         2.0              0         0.0      0.0
2         1   48         1.0              1        20.0      0.0
3         0   61         3.0              1        30.0      0.0
4         0   46         3.0              1        23.0      0.0
...    ...  ...  ...          ...          ...      ...
4233      1   50         1.0              1         1.0      0.0
4234      1   51         3.0              1        43.0      0.0
4235      0   48         2.0              1        20.0      NaN
4236      0   44         1.0              1        15.0      0.0
4237      0   52         2.0              0         0.0      0.0

      prevalentStroke  prevalentHyp  diabetes  totChol  sysBP  diaBP  BMI  \
0                   0              0         0    195.0   106.0   70.0  26.97
1                   0              0         0    250.0   121.0   81.0  28.73
2                   0              0         0    245.0   127.5   80.0  25.34
3                   0              1         0    225.0   150.0   95.0  28.58
4                   0              0         0    285.0   130.0   84.0  23.10
...    ...          ...          ...    ...    ...    ...    ...
4233      0          1         0    313.0   179.0   92.0  25.97
4234      0          0         0    207.0   126.5   80.0  19.71
4235      0          0         0    248.0   131.0   72.0  22.00
4236      0          0         0    210.0   126.5   87.0  19.16
4237      0          0         0    269.0   133.5   83.0  21.47

      heartRate  glucose  TenYearCHD
0         80.0     77.0           0
1         95.0     76.0           0
2         75.0     70.0           0
3         65.0    103.0           1

```

4	85.0	85.0	0
...
4233	66.0	86.0	1
4234	65.0	68.0	0
4235	84.0	86.0	0
4236	86.0	NaN	0
4237	80.0	107.0	0

[4238 rows x 16 columns]

1 MISSING VALUE TREATMENT

Since, 'glucose' and 'education' columns had a significant amount of null values, so we replaced them with the mean of values for their respective columns

```
[27]: df['glucose'].fillna(value = df['glucose'].mean(),inplace=True)
```

```
[29]: df['education'].fillna(value = df['education'].mean(),inplace=True)
```

```
[31]: df['heartRate'].fillna(value = df['heartRate'].mean(),inplace=True)
```

```
[33]: df['BMI'].fillna(value = df['BMI'].mean(),inplace=True)
```

```
[35]: df['cigsPerDay'].fillna(value = df['cigsPerDay'].mean(),inplace=True)
```

```
[37]: df['totChol'].fillna(value = df['totChol'].mean(),inplace=True)
```

```
[39]: df['BPMeds'].fillna(value = df['BPMeds'].mean(),inplace=True)
```

```
[41]: df.isna().sum()
```

```
[41]: male          0
      age          0
      education    0
      currentSmoker 0
      cigsPerDay    0
      BPMeds        0
      prevalentStroke 0
      prevalentHyp  0
      diabetes      0
      totChol       0
      sysBP         0
      diaBP         0
      BMI           0
      heartRate     0
      glucose       0
```

```
TenYearCHD      0
dtype: int64
```

```
[43]: #Splitting the dependent and independent variables.
```

```
x = df.drop("TenYearCHD",axis=1)
y = df['TenYearCHD']
```

```
[45]: x #checking the features
```

```
[45]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	\
0	1	39	4.0	0	0.0	0.00000	
1	0	46	2.0	0	0.0	0.00000	
2	1	48	1.0	1	20.0	0.00000	
3	0	61	3.0	1	30.0	0.00000	
4	0	46	3.0	1	23.0	0.00000	
...	
4233	1	50	1.0	1	1.0	0.00000	
4234	1	51	3.0	1	43.0	0.00000	
4235	0	48	2.0	1	20.0	0.02963	
4236	0	44	1.0	1	15.0	0.00000	
4237	0	52	2.0	0	0.0	0.00000	

	prevalentStroke	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	\
0	0	0	0	195.0	106.0	70.0	26.97	
1	0	0	0	250.0	121.0	81.0	28.73	
2	0	0	0	245.0	127.5	80.0	25.34	
3	0	1	0	225.0	150.0	95.0	28.58	
4	0	0	0	285.0	130.0	84.0	23.10	
...	
4233	0	1	0	313.0	179.0	92.0	25.97	
4234	0	0	0	207.0	126.5	80.0	19.71	
4235	0	0	0	248.0	131.0	72.0	22.00	
4236	0	0	0	210.0	126.5	87.0	19.16	
4237	0	0	0	269.0	133.5	83.0	21.47	

	heartRate	glucose
0	80.0	77.000000
1	95.0	76.000000
2	75.0	70.000000
3	65.0	103.000000
4	85.0	85.000000
...
4233	66.0	86.000000
4234	65.0	68.000000
4235	84.0	86.000000
4236	86.0	81.966753
4237	80.0	107.000000

[4238 rows x 15 columns]

2 TRAIN TEST SPLIT

```
[48]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.  
↪2,random_state=42)
```

```
[50]: y_train
```

```
[50]: 3252    0  
      3946    0  
      1261    0  
      2536    0  
      4089    0  
      ..  
      3444    0  
      466     0  
      3092    0  
      3772    0  
      860     0  
      Name: TenYearCHD, Length: 3390, dtype: int64
```

3 Logistic Regression Algorithm

```
[53]: from sklearn.linear_model import LogisticRegression  
      model = LogisticRegression().fit(x_train,y_train)  
      model.score(x_train, y_train)
```

```
[53]: 0.8489675516224189
```

```
[ ]:
```